

Outline

- The Planning Problem
- Planning with State-Based Search
- Partial-Order Planning
- Planning with Propositional Logic

Example: Single-Player "Game"

Legality

```
legal(you, rightShoe) <= true(rightSockOn)
legal(you, rightSock)
legal(you, leftShoe) <= true(leftSockOn)
legal(you, leftSock)
```

Update

```
next(rightShoeOn) <= does(you, rightShoe)
next(rightSockOn) <= does(you, rightSock)
next(leftShoeOn) <= does(you, leftShoe)
next(leftSockOn) <= does(you, leftSock)
```

Termination and Goal

```
terminal <= true(rightShoeOn) ^ true(leftShoeOn)
goal(you, 100) <= true(rightShoeOn) ^ true(leftShoeOn)
```

Applications of Planning



A Simpler Description Language for Planning Problems

- Initial state: conjunction of variable-free atoms
- Actions: <Name, Precondition, Effect>
 - Name: Action name + parameter list
 - Precond: Conjunction of literals
 - Effect: Conjunction of literals
- Goal: logical sentence

A solution to a planning problem is an action sequence that, when executed in the initial state, results in a state that satisfies the goal.

Example

- Initial state
()
- Actions


rightShoe	Precond:	rightSockOn
	Effect:	rightShoeOn
rightSock	Effect:	rightSockOn
leftShoe	Precond:	leftSockOn
	Effect:	leftShoeOn
leftSock	Effect:	leftSockOn
- Goal
 $\text{rightShoeOn} \wedge \text{leftShoeOn}$

... Formalised in the Planning Description Language

- Initial state
 $\text{on}(a, \text{table}) \wedge \text{on}(b, \text{table}) \wedge \text{on}(c, a) \wedge \text{clear}(b) \wedge \text{clear}(c)$
- Actions

Name:	$\text{move}(X, Y, Z)$
Precond:	$\text{on}(X, Y) \wedge \text{clear}(X) \wedge \text{clear}(Z) \wedge X \neq Z \wedge Y \neq Z$
Effect:	$\text{on}(X, Z) \wedge \text{clear}(Y) \wedge \neg \text{on}(X, Y) \wedge \neg \text{clear}(Z)$
Name:	$\text{moveToTable}(X, Y)$
Precond:	$\text{on}(X, Y) \wedge \text{clear}(X)$
Effect:	$\text{on}(X, \text{table}) \wedge \text{clear}(Y) \wedge \neg \text{on}(X, Y)$
- Goal
 $\text{on}(a, b) \wedge \text{on}(b, c)$


Another Example: Blocks World Planning




A robot arm can pick up a block and move it to another position.
The arm can only pick up one block at a time.

Planning by State-Based Search

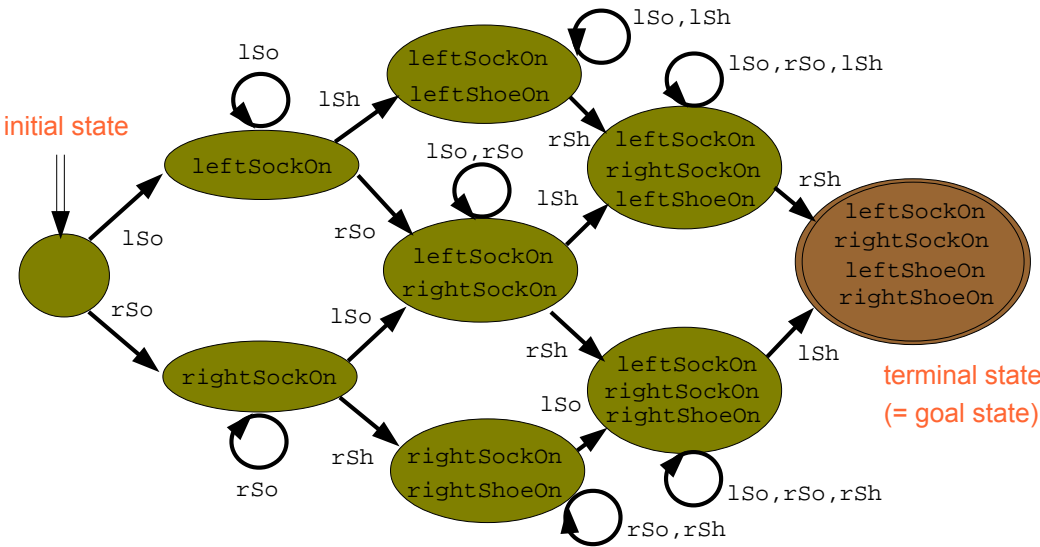
Recap: State Machines




Forward Search




State Machine for the Example Game



Backward Search

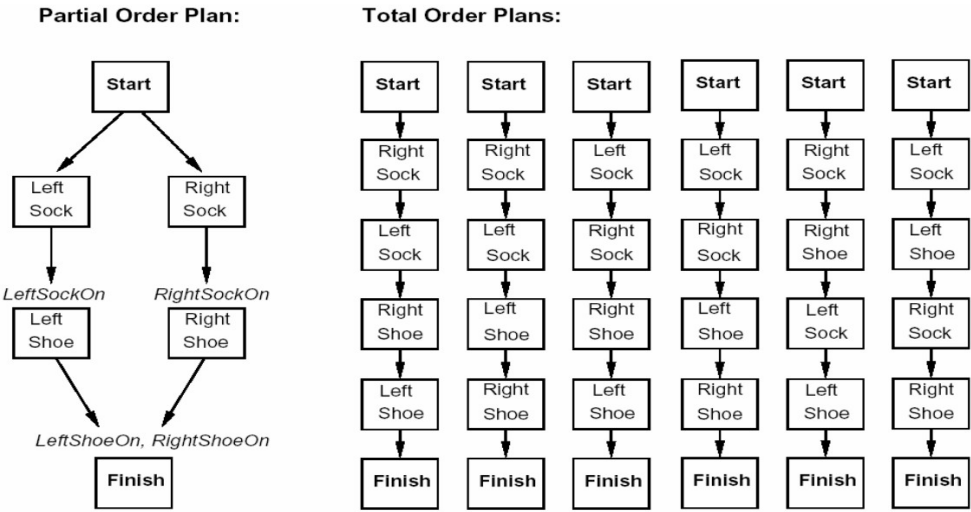


Bidirectional Search

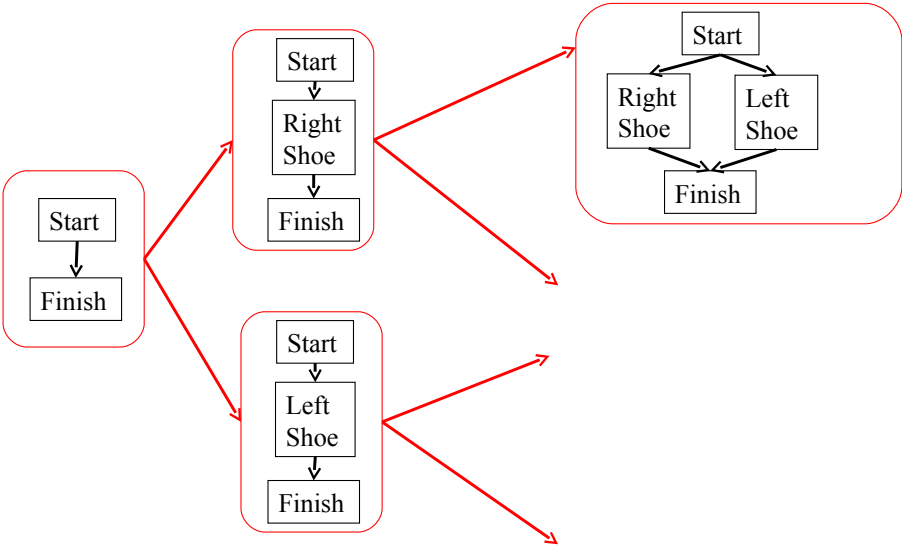


Partial-Order Planning

Partial-Order Plan (POP): Example



Plan-Space Search with POPs



Partial-Order Planning as Search Problem

- **Search nodes** are (mostly unfinished) partial-order plans
The initial plan contains only *Start* and *Finish* action
- **Plans** have 4 components:
 - A set of actions (steps of the plan)
 - A set of ordering constraints $A < B$ (A before B)
 - A set of causal links $A \xrightarrow{p} B$ (read: "A achieves p for B")
 - A set of open preconditions
- A plan is **consistent** if there are no cycles in the ordering constraints and no conflicts with the causal links.
- An action C **conflicts** with a causal link $A \xrightarrow{p} B$ if C has the effect $\neg p$ and C could come after A and before B
- A consistent plan with no open preconditions is a **solution**.

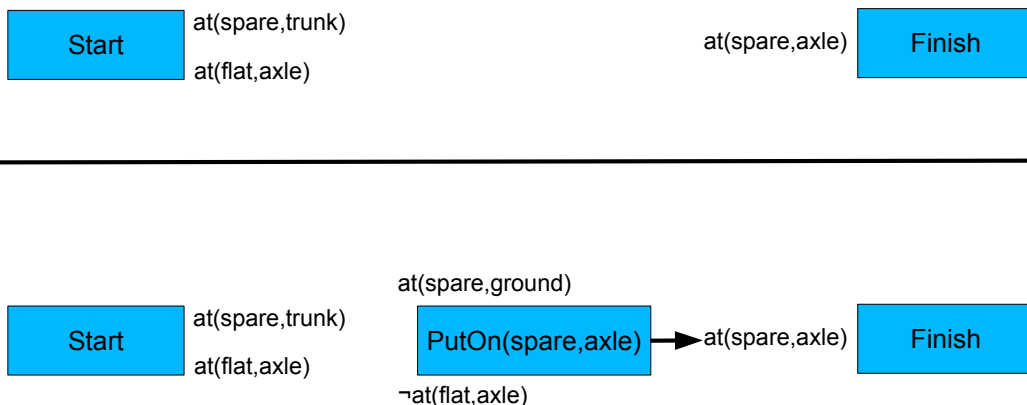
Example: Flat Tire Problem

- **Initial state** `at(flat, axle) ∧ at(spare, trunk)`
- **Actions**
 - Name: `remove(spare, trunk)`
Precond: `at(spare, trunk)`
Effect: `¬at(spare, trunk) ∧ at(spare, ground)`
 - Name: `remove(flat, axle)`
Precond: `at(flat, axle)`
Effect: `¬at(flat, axle) ∧ at(flat, ground)`
 - Name: `putOn(spare, axle)`
Precond: `at(spare, ground) ∧ ¬at(flat, axle)`
Effect: `¬at(spare, ground) ∧ at(spare, axle)`
- **Goal** `at(spare, axle)`

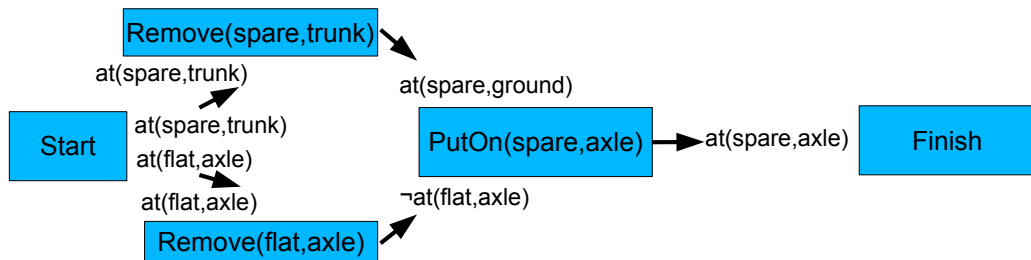

Algorithm for Solving POPs

- Define effect of *Start* := initial state of the planning problem (no precondition)
- Define precondition of *Finish* := goal of the planning problem (no effect)
- The initial plan contains *Start* and *Finish*, the ordering constraint $Start < Finish$, no causal links. All preconditions of *Finish* are open.
- **Repeat**
 - Pick an open precondition p (of an action B in the plan)
 - Pick an action A with effect p
 - Add the causal link $A \xrightarrow{p} B$ and the ordering constraint $A < B$ (if A is new to the plan, add $Start < A$ and $A < Finish$)
 - If a conflict arises between a causal link $A \xrightarrow{p} B$ and an action C: add either $B < C$ or $C < A$
- **Retry** (with different choices) if plan is inconsistent, stop if solution is found

POP for the Flat Tire Problem (1)



POP for the Flat Tire Problem (2)




Encoding Planning Problems in Propositional Logic

- Planning can be done by testing the **satisfiability** of a logical sentence:

$$\text{initial state} \wedge \text{all possible actions} \wedge \text{goal}$$
- This sentence contains propositions for every action occurrence
 - A model will assign *true* to an action A iff A is part of the correct plan
- An assignment that corresponds to an incorrect plan will not be a model because of inconsistency with the assertion that *goal* is true
- If the planning problem is unsolvable, there will be no model for the sentence
- Planners based on satisfiability can handle large planning problems

Planning with Propositional Logic

Recap: Blocks World Planning



A robot arm can pick up a block and move it to another position.
 The arm can only pick up one block at a time.

Example: Blocks World Planning as Satisfiability (1)

- Encoding of the initial state


```
on(a,table)^0
on(b,table)^0
on(c,a)^0
clear(b)^0
clear(c)^0
```
- Encoding of action preconditions


```
move(X,Y,Z)^T => on(X,Y)^T ^ clear(X)^T ^ clear(Z)^T
moveToTable(X,Y)^T => on(X,Y)^T ^ clear(X)^T
            (for all X,Y,Z∈{a,b,c,table}, T∈{0,1,2,...,max-1}, X≠Z, Y≠Z)
```
- Action exclusion axioms


```
¬(move(X,Y,Z)^T ^ moveToTable(X',Y')^T)
¬(moveToTable(X,Y)^T ^ moveToTable(X',Y')^T)
¬(move(X,Y,Z)^T ^ move(X',Y',Z')^T)          (for suitable X,X',...)
```

Example: Blocks World Planning as Satisfiability (2)

- Encoding of action effects


```
move(X,Y,Z)^T => on(X,Z)^T+1 ^ clear(Y)^T+1
move(X,Y,Z)^T => ¬on(X,Y)^T+1 ^ ¬clear(Z)^T+1
moveToTable(X,Y)^T => on(X,table)^T+1 ^ clear(Y)^T+1
moveToTable(X,Y)^T => ¬on(X,Y)^T+1
```
- Explanation closure axioms


```
on(X,Z)^T+1 => on(X,Z)^T ∨ move(X,Y,Z)^T ∨
            (moveToTable(X,Y)^T ^ Z=table)
clear(Y)^T+1=> clear(Y)^T ∨
            move(X,Y,Z)^T ∨ moveToTable(X,Y)^T
```
- Encoding of the goal



```
on(a,b)^max ^ on(b,c)^max
```

Solution (max=3): a model that contains

`moveTable(c,a)^0, move(b,table,c)^1, move(a,table,b)^2`


Conditional Planning

Planning Under Incomplete Information: Maze World




Initial State: \textcircled{ac} (robot in a, gold in c)


Environment Model




Agent Percepts



Agent Actions



Initial State and Goal



Planning

Planning is the process of finding a transition diagram *for our agent* that causes its environment to go from any initial state to a goal state.



Planning can be done *offline* and the resulting plan/program installed in the agent *or* the planning can be done *online* followed by execution.

Incompleteness

Possible sources of incompleteness:

Partial knowledge of

- Initial state
- Transition diagram for environment
- Goal


Complete Planning Techniques under incomplete information

- Coercion (e.g. do the *grab* action at all locations)
- Conditional plan (e.g. if see the gold grab it; else move)


Postponement Techniques

- Delayed planning


State Space Planning



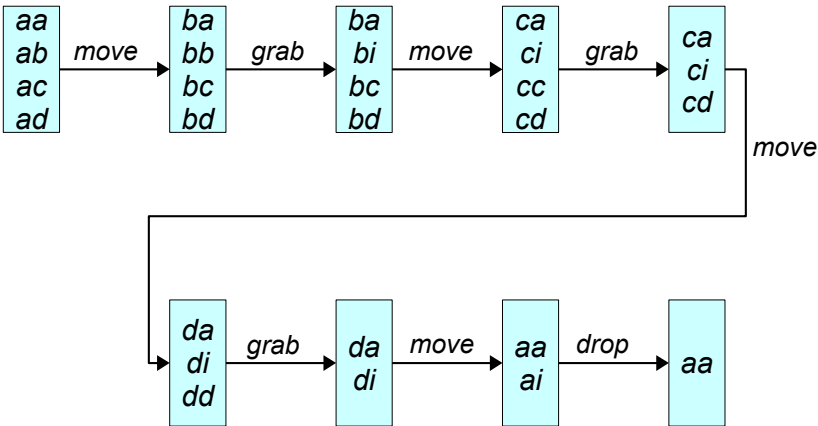
Initial State Uncertainty




Sequential State Set Progression




Sequential State Set Plan




Plan Execution



Conditional State Set Progression



Conditional State Set Plan



Background Reading

Planning

- Russell & Norvig AIMA (3rd ed): Chapter 10
(2nd edition: Chapter 11)

Comparison

Sequential plan

- possible that no plan exists
- plan may contain redundant moves

Conditional plan

- large search space

Delayed planning

- irreversibility problematic

As we can see from this analysis, it is sometimes desirable for an agent to do only a portion of its planning up front, secure in the knowledge that it can do more later as necessary.

Planning can be done **offline** and the resulting plan/program executed during play *or* the planning can be done **online** and interleaved with execution.